

# Project 2

*Kyle Bartsch*

*November 5, 2015*

## Summary

This dataset was pulled from the UC Irvine Machine Learning Repository:

<http://archive.ics.uci.edu/ml/index.html> (<http://archive.ics.uci.edu/ml/index.html>). It is compiled from data related to direct marketing campaigns of a Portuguese banking institution. The marketing campaigns were based on phone calls....The classification goal is to predict if the client will subscribe to a term deposit.

## Data Import and Coding

### Dataset Attributes:

The data consists of 45211 observations of 21 variables. Half of the data was withheld as a test dataset for model validation.

#### Input variables:

# bank client data:

1 - age (numeric)

2 - job : type of job (categorical: 'admin.', 'blue-collar', 'entrepreneur', 'housemaid', 'management', 'retired', 'self-employed', 'services', 'student', 'technician', 'unemployed', 'unknown')

3 - marital : marital status (categorical: 'divorced', 'married', 'single', 'unknown'; note: 'divorced' means divorced or widowed)

4 - education (categorical:

'basic.4y', 'basic.6y', 'basic.9y', 'high.school', 'illiterate', 'professional.course', 'university.degree', 'unknown')

5 - default: has credit in default? (categorical: 'no', 'yes', 'unknown')

6 - housing: has housing loan? (categorical: 'no', 'yes', 'unknown')

7 - loan: has personal loan? (categorical: 'no', 'yes', 'unknown')

# related with the last contact of the current campaign:

8 - contact: contact communication type (categorical: 'cellular', 'telephone')

9 - month: last contact month of year (categorical: 'jan', 'feb', 'mar', ..., 'nov', 'dec')

10 - day\_of\_week: last contact day of the week (categorical: 'mon', 'tue', 'wed', 'thu', 'fri')

11 - duration: last contact duration, in seconds (numeric). Important note: this attribute highly affects the output target (e.g., if duration=0 then y='no'). Yet, the duration is not known before a call is performed.

Also, after the end of the call y is obviously known. Thus, this input should only be included for benchmark purposes and should be discarded if the intention is to have a realistic predictive model.

# other attributes:

12 - campaign: number of contacts performed during this campaign and for this client (numeric, includes last contact)

13 - pdays: number of days that passed by after the client was last contacted from a previous campaign (numeric; 999 means client was not previously contacted)

14 - previous: number of contacts performed before this campaign and for this client (numeric)

15 - poutcome: outcome of the previous marketing campaign (categorical: 'failure', 'nonexistent', 'success')

# social and economic context attributes

16 - emp.var.rate: employment variation rate - quarterly indicator (numeric)

17 - cons.price.idx: consumer price index - monthly indicator (numeric)

18 - cons.conf.idx: consumer confidence index - monthly indicator (numeric)

19 - euribor3m: euribor 3 month rate - daily indicator (numeric)

20 - nr.employed: number of employees - quarterly indicator (numeric)

Output variable (desired target):

21 - y - has the client subscribed a term deposit? (binary: 'yes', 'no')

The online directions state that the "Duration" variable highly affects the output target and should be only used as a benchmark; it should be removed if the intention is to develop a truly predictive model:

```
d = d[,-which(names(d) %in% c("X","Duration"))] #remove index and duration
summary(d)
```

```

##           Age           Job           Marital
## Min.      :17.0    admin.      :5189    divorced: 2285
## 1st Qu.:32.0    blue-collar:4592    married :12500
## Median :38.0    technician :3417    single  : 5767
## Mean      :40.1    services   :1967    unknown :   41
## 3rd Qu.:47.0    management :1490
## Max.      :98.0    retired    : 874
##           (Other)    :3064
##           Education    Default           Housing
## university.degree :6091    no      :16317    no      : 9346
## high.school        :4777    unknown: 4275    unknown: 497
## basic.9y           :3001    yes      :    1    yes      :10750
## professional.course:2660
## basic.4y           :2073
## basic.6y           :1117
## (Other)            : 874
##           Loan           Contact           Month           Day
## no      :17011    cellular :13157    may      :6907    fri:3905
## unknown: 497    telephone: 7436    jul      :3590    mon:4178
## yes      : 3085                                aug      :3148    thu:4317
##                                           jun      :2600    tue:4102
##                                           nov      :2046    wed:4091
##                                           apr      :1306
##                                           (Other): 996
##           Campaign           pDays           Previous           pOutcome
## Min.      : 1.000    Min.      : 0.0    Min.      :0.000    failure    : 2090
## 1st Qu.: 1.000    1st Qu.:999.0    1st Qu.:0.000    nonexistent:17822
## Median : 2.000    Median :999.0    Median :0.000    success    : 681
## Mean      : 2.563    Mean      :962.7    Mean      :0.173
## 3rd Qu.: 3.000    3rd Qu.:999.0    3rd Qu.:0.000
## Max.      :43.000    Max.      :999.0    Max.      :7.000
##
##           EmpVarRate           ConsPriceIdx           ConsConfIdx           Euribor3M
## Min.      :-3.40000    Min.      :92.20    Min.      :-50.80    Min.      :0.634
## 1st Qu.: -1.80000    1st Qu.:93.08    1st Qu.: -42.70    1st Qu.:1.344
## Median : 1.10000    Median :93.75    Median : -41.80    Median :4.857
## Mean      : 0.08379    Mean      :93.57    Mean      : -40.47    Mean      :3.623
## 3rd Qu.: 1.40000    3rd Qu.:93.99    3rd Qu.: -36.40    3rd Qu.:4.961
## Max.      : 1.40000    Max.      :94.77    Max.      : -26.90    Max.      :5.045
##
##           NREmployed           Subscribed
## Min.      :4964    no :18228
## 1st Qu.:5099    yes: 2365
## Median :5191
## Mean      :5167
## 3rd Qu.:5228
## Max.      :5228
##

```

```
str(d)
```

```
## 'data.frame':    20593 obs. of  20 variables:
## $ Age           : int  46 28 44 38 34 31 33 44 40 40 ...
## $ Job           : Factor w/ 12 levels "admin.,"blue-collar",...: 1 8 10 10 10 1 1 11 1
10 ...
## $ Marital       : Factor w/ 4 levels "divorced","married",...: 1 3 2 2 2 2 2 2 2 ...
## $ Education     : Factor w/ 8 levels "basic.4y","basic.6y",...: 7 4 6 7 6 7 3 4 7 7 ...
## $ Default       : Factor w/ 3 levels "no","unknown",...: 2 1 2 1 1 1 1 2 2 1 ...
## $ Housing       : Factor w/ 3 levels "no","unknown",...: 1 3 3 2 3 1 2 3 3 1 ...
## $ Loan          : Factor w/ 3 levels "no","unknown",...: 1 3 1 2 1 1 2 1 3 1 ...
## $ Contact       : Factor w/ 2 levels "cellular","telephone": 1 1 2 1 1 2 2 1 2 1 ...
## $ Month         : Factor w/ 10 levels "apr","aug","dec",...: 4 7 5 7 2 4 7 2 7 8 ...
## $ Day           : Factor w/ 5 levels "fri","mon","thu",...: 3 5 1 3 1 2 4 5 5 1 ...
## $ Campaign      : int   1 1 2 1 1 1 1 1 5 1 ...
## $ pDays         : int  999 999 999 999 999 999 999 999 999 999 ...
## $ Previous      : int   0 0 0 1 0 0 0 0 0 1 ...
## $ pOutcome      : Factor w/ 3 levels "failure","nonexistent",...: 2 2 2 1 2 2 2 2 2 1
...
## $ EmpVarRate    : num   1.4 -1.8 1.4 -1.8 1.4 -1.7 1.1 1.4 1.1 -0.1 ...
## $ ConsPriceIdx  : num   93.9 92.9 94.5 92.9 93.4 ...
## $ ConsConfIdx   : num  -42.7 -46.2 -41.8 -46.2 -36.1 -40.3 -36.4 -36.1 -36.4 -42 ...
## $ Euribor3M     : num   4.96 1.28 4.97 1.33 4.97 ...
## $ NREmployed    : num  5228 5099 5228 5099 5228 ...
## $ Subscribed    : Factor w/ 2 levels "no","yes": 1 2 1 1 1 1 1 1 1 1 ...
```

The data appears to be accurately coded: continuous variables are coded as integers or numbers and categorical variables are coded as factors.

## Missing Values

Let's see look at which variables are missing. For this dataset, missing values have been coded as "unknown":

```
m = matrix(data=NA,nrow=length(names(d)),ncol=2)
dimnames(m) = list(names(d),c("Total Missing","Percent Missing"))
for(i in 1:nrow(m)){
  l = length(d[,i])
  u = length(which(d[,i]=="unknown"))
  m[i,"Total Missing"] = u
  m[i,"Percent Missing"] = round(u/l,digits=4)
}
m[m[, "Percent Missing"]!=0,]
```



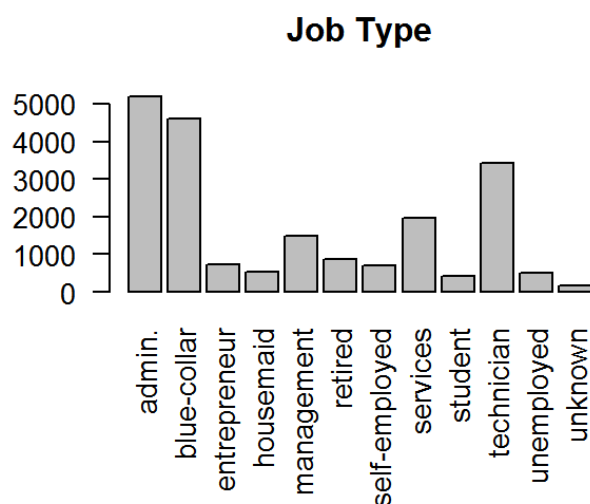

##	Total	Missing	Percent Missing
## Job	162		0.0079
## Marital	41		0.0020
## Education	865		0.0420
## Default	4275		0.2076
## Housing	497		0.0241
## Loan	497		0.0241

Here we see that we have 6 variables with missing values, ranging from under 1% missing up to just over 20% missing. “Default” is the most problematic variable here, with 20% of observations missing. At this point, we could try and impute missing values using techniques like tree-based methods or K-Nearest Neighbors. However, in the interest of exploring how different models can handle missing values, we will not attempt to make any imputations.


## Exploratory Data Analysis

Now that we have updated incorrect datatypes and recoded missing data, let’s do some exploratory data analysis by looking at univariate plots of the data.


### Univariate Plots




**Credit Default Status**




**Has Housing Loan?**




**Has Personal Loan?**




**Contact Method**




**Last Contact Month of Year**




**Last Contact Day of Week**




**Campaign**



**pDays**



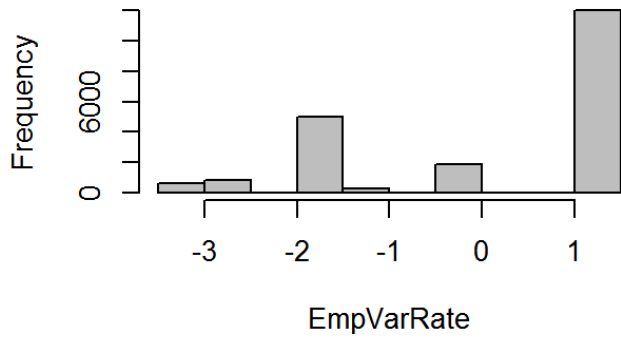
**Previous**



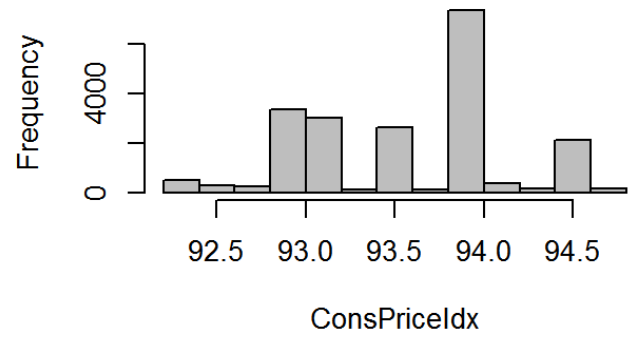
**Outcome of Prior Campaign**



**Employment Variation Rate**




**Consumer Price Index**

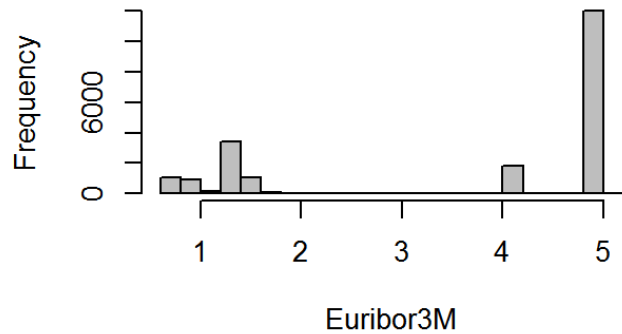




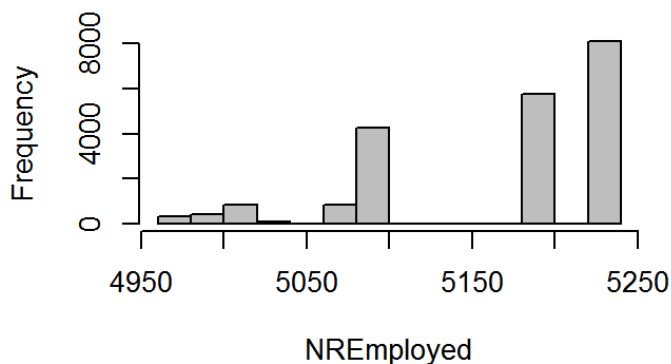
**Consumer Confidence Index**




**Euribor 3 Month Rate**



**Number of Employees**



**Did the Client Subscribe?**




A few generalized conclusions from looking at univariate plots:

- Several variables are quite skewed/disproportionate: Default, Personal Loan, Campaign, pDays, Previous
- Most of these clients were not previously contacted
- pDays needs to be addressed. The instructions state that "if the client was not previously contacted, the variable is coded as"999". We will recode this variable and make it more simple by splitting it into categories based on weeks since last contact:

```
d$pContact = as.factor(ifelse(d$pDays<=7,"1 Week",ifelse(d$pDays<=14,"2 Weeks",ifelse(d$pDays<=21,"3 Weeks",ifelse(d$pDays<=28,"4 Weeks","No Contact")))))  
d=d[, -which(names(d) %in% c("pDays"))] #remove pDays variable  
barplot(table(d$pContact),main="Weeks Passed Prev Campaign",las=2)
```


## Weeks Passed Prev Campaign




## Bivariate Plots

Now let's examine how each of our variables relates to the subscription rate by examining bivariate plots. NOTE: for each of the mosaic plots below (categorical predictors), a dotted red line has been included that marks the proportion of yes/no for the response accross all observations (~88.52%).

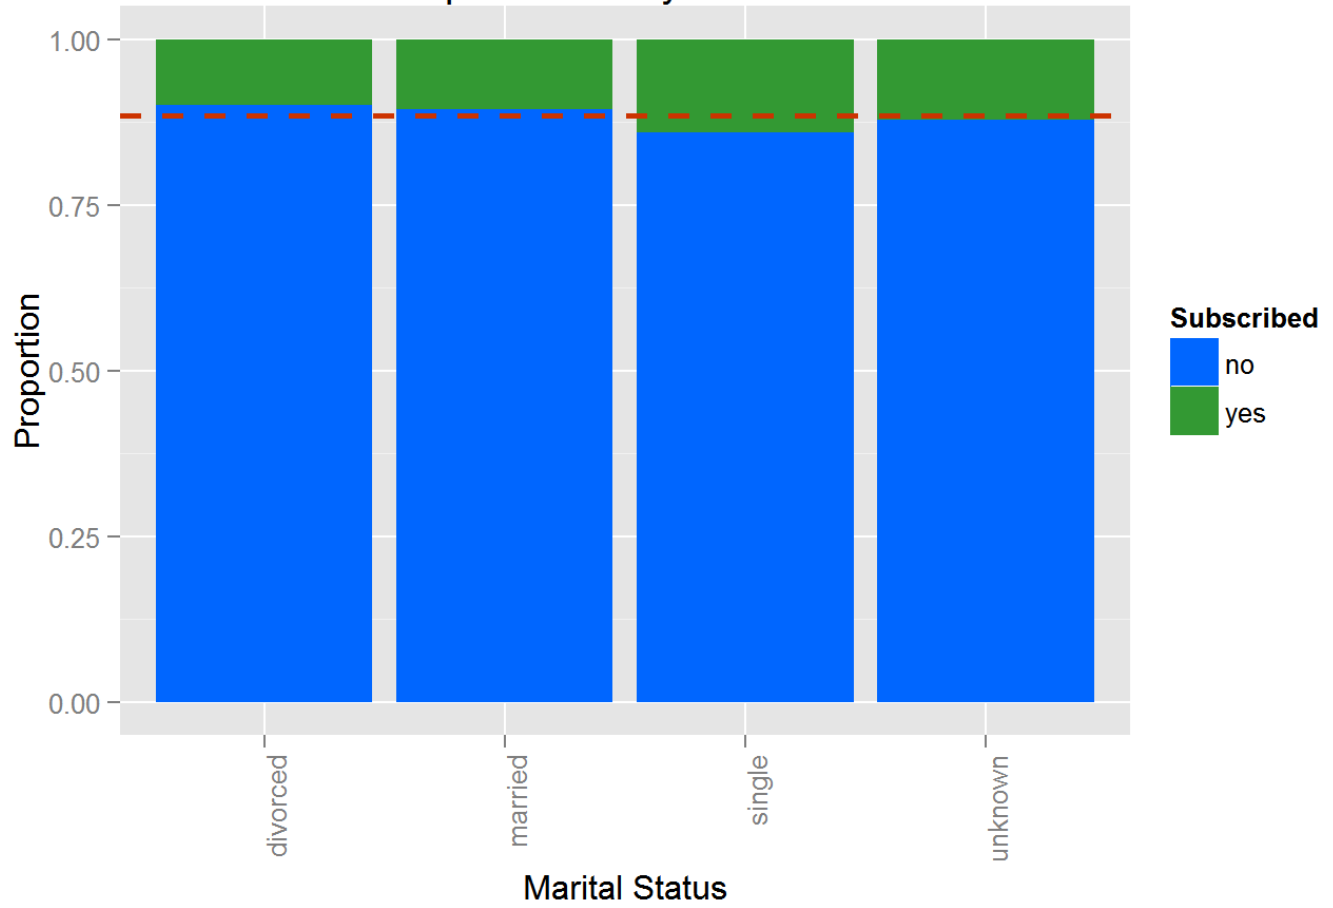
Age vs. Subscription




Subscription Rate by Job Type



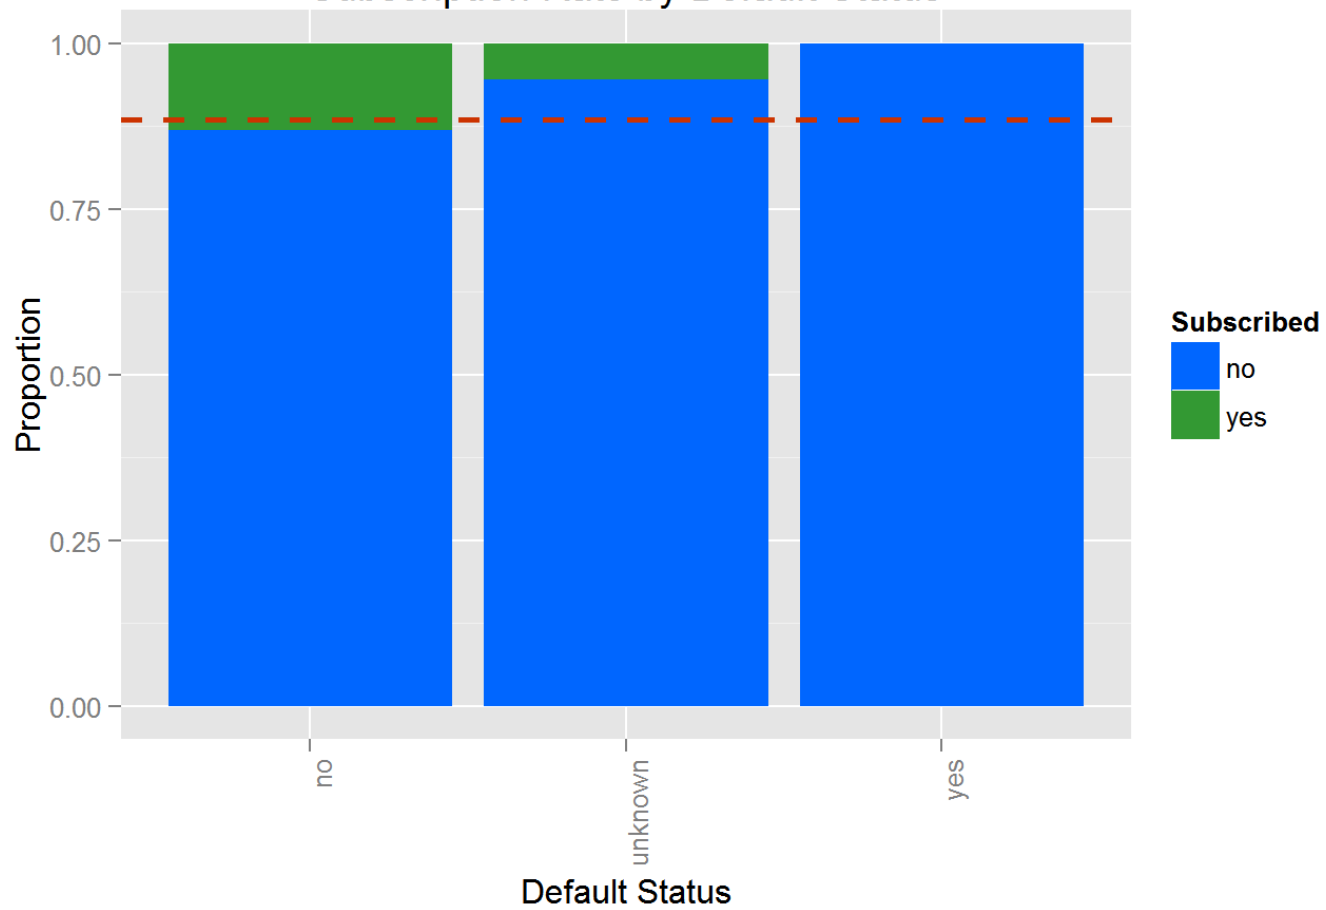
Subscription Rate by Marital Status




Subscription Rate by Education



Subscription Rate by Default Status




Subscription Rate by Housing Loan Status

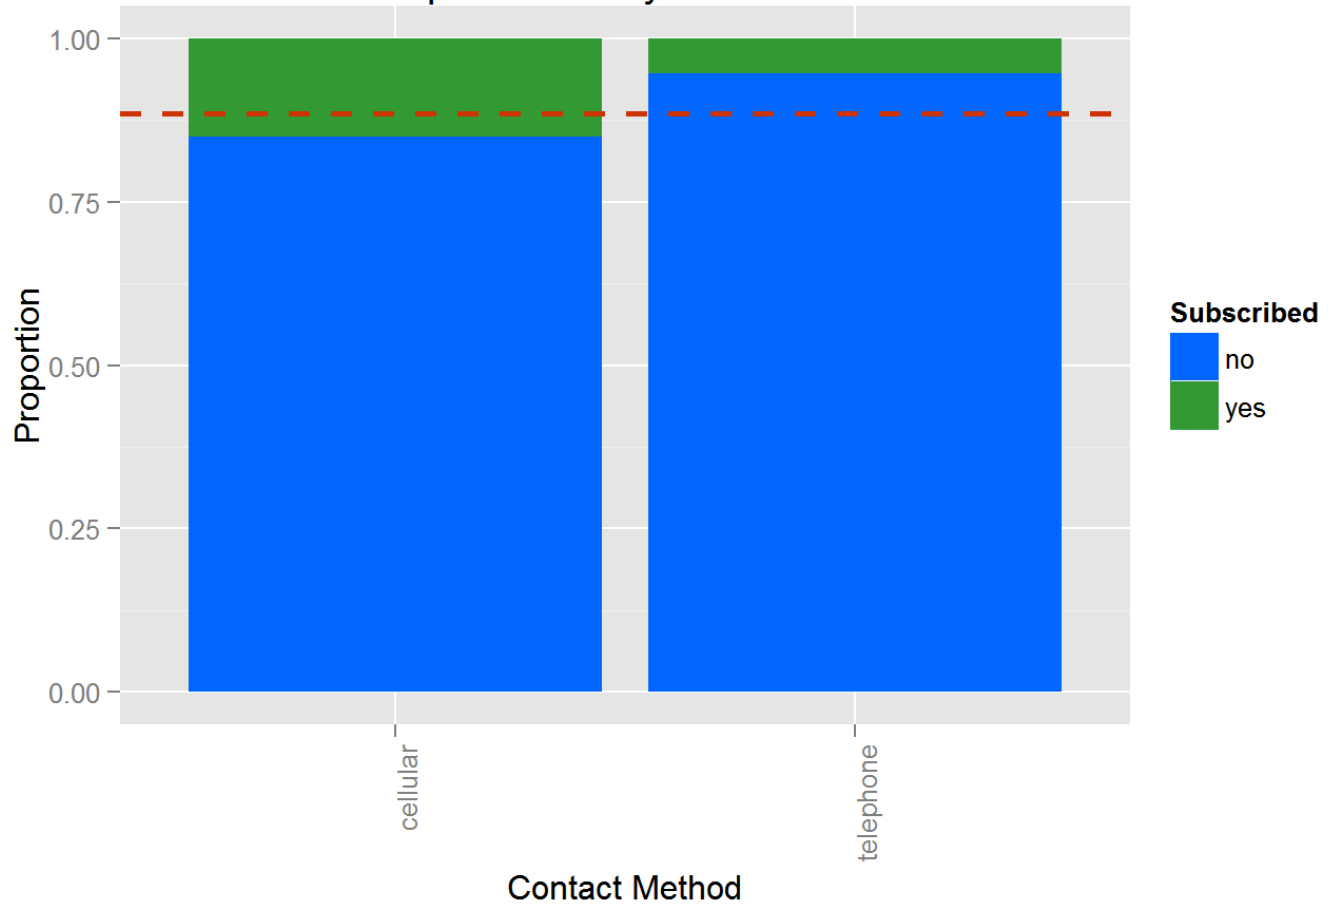





Subscription Rate by Person Loan Status




Subscription Rate by Contact Method




Subscription Rate by Month



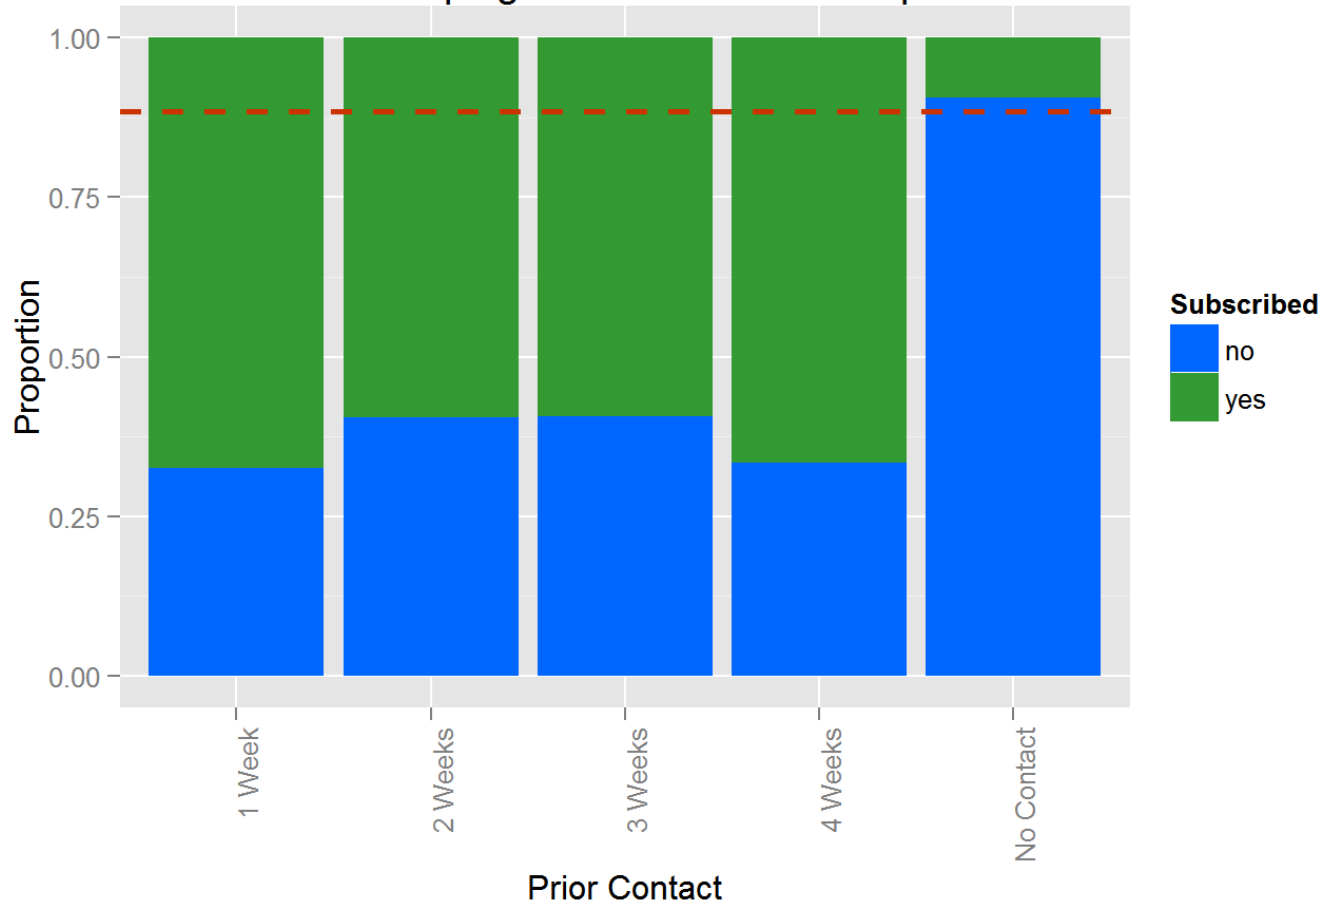
Subscription Rate by Day of Week




Campaign vs. Subscription




Prior Campaign Contact vs. Subscription



Previous Contacts Prior to Campaign vs. Subscription




Outcome of Prior Campaign vs. Subscription







Employment Variation Rate vs. Subscription




Consumer Price Index vs. Subscription




Consumer Confidence Index vs. Subscription



Euribor 3 Month Rate vs. Subscription





#### A few generalized conclusions from looking at bivariate plots:

The following variables exhibit variation in the response and may be strong predictors:

- Job - Retired and Student job types had higher rates of subscription compared to other job types.
- Default Status - Clients who have defaulted did not subscribe at all.
- Contact Method - Clients who were contacted by cell phone subscribed more than those by telephone.
- Month - March, April, September October and December have much higher rates of subscription than others.
- Prior Contact - Clients that had contact of any kind had much higher rates of subscription than those that were never contacted.
- Pervious - Clients that had been contacted in the previous campaign subscribed at a higher rate.
- Prior Outcome - Clients that subscribed in the prior campaign subscribed in this campaign at a higher rate.
- Euibor 3 Month Rate - The lower the rate, it appears the subscription rate increases.
- Number of Employees - Subscription rates seem to be higher at lower employment levels.

## Modeling

There are several different ways we can model a binary response variable. The classic model is Logistic Regression. For this project, we will be going further and also applying Classification Trees, Bagging, Random Forests and Gradient Boosting. The goal of using five different models is to compare how each

model performs on the test dataset we have withheld.

## Assumptions

For scenarios that involve classification, there needs to be special consideration to the costs of False Positives vs. False Negatives. In this case, we intend to use the results of our models to send a direct mailer to potential bank customers. For a real-world scenario, we would have to determine the costs and benefits of producing a direct mailer, gaining a customer, not gaining a customer, increasing direct competition, etc. For our sake, we will assume that the cost of generating the mailer is much less than that of not gaining new customers. Further, if we assume that a new customer would generate revenue well beyond the cost of the mailer, and that customers not receiving the mailer will not consider our bank, then we can assume that False Negatives (not sending a mailer when they would have responded) are more costly than False Positives (sending the mailer, not having the customer respond). Additionally, we can assume that unless we reach out to a customer, perhaps they would choose a different bank, which would result in an increase in direct competition. Again, here we would want to minimize False Negatives. Therefore, for all modeling, we will set the cost of False Negatives as **2x** that of False Positives.

## Train Set and Validation Set

For modeling, we will split our dataset into 80% for training and 20% for validation:

```
set.seed(72) #set seed for reproducibility
train.ind = sample(nrow(d), nrow(d)/5) #split into 5ths
train = d[-train.ind,] #assign training to 80%
val = d[train.ind,] #assign validation to 20%
uc.val = c(0,1)[unclass(val$Subscribed)] #create an unclassed vector of responses from val set
```


## Logistic Regression

We already have a good idea of which variables might be beneficial in a model for this problem. However, we have to be wary of overfitting the model and should look at which variables might be the most predictive. Here we use forward selection to build the best models based on the AIC criterion:

```
full.model = glm(Subscribed~.-Subscribed,data=train,family=binomial(link="logit"))
null.model = glm(Subscribed~1,data=train,family=binomial(link="logit"))
step.model = step(null.model,scope=list(upper=full.model),data=train,direction="both")
```

##		Step	Df	Deviance	Resid. Df	Resid. Dev	AIC
## 1		NA	NA	16474	11714.546	11716.546	
## 2	+	NREmployed	-1	1867.123982	16473	9847.422	9851.422
## 3	+	pOutcome	-2	308.434291	16471	9538.988	9546.988
## 4	+	Month	-9	236.612471	16462	9302.375	9328.375
## 5	+	Contact	-1	36.612642	16461	9265.763	9293.763
## 6	+	Day	-4	24.749391	16457	9241.013	9277.013
## 7	+	pContact	-4	20.921131	16453	9220.092	9264.092
## 8	+	Campaign	-1	13.514836	16452	9206.577	9252.577
## 9	+	ConsConfIndx	-1	9.764857	16451	9196.812	9244.812
## 10	+	Default	-1	7.171042	16450	9189.641	9239.641
## 11	+	Previous	-1	2.760272	16449	9186.881	9238.881
## 12	+	Age	-1	2.407898	16448	9184.473	9238.473

Stepwise AIC by Total Variables Added



Looks like we hit the minimum AIC around 6 or 7 variables added. There is little benefit to including more variables as we would increase the risk in overfitting the model. Let's build a Logistic Model using the first 7 variables from the stepwise method:


```
log.mod = glm(Subscribed~NREmployed+pOutcome+Month+Contact+Day+pContact+Campaign,
              data=train, family = binomial(link="logit"))
summary(log.mod)
```

```
##
## Call:
## glm(formula = Subscribed ~ NREmployed + pOutcome + Month + Contact +
##      Day + pContact + Campaign, family = binomial(link = "logit"),
##      data = train)
##
## Deviance Residuals:
##      Min        1Q    Median        3Q        Max
## -2.0662  -0.3936  -0.3306  -0.2468   3.0124
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    58.923495   2.273638  25.916 < 2e-16 ***
## NREmployed     -0.011540   0.000452 -25.529 < 2e-16 ***
## pOutcomenonexistent  0.460486   0.087396   5.269 1.37e-07 ***
## pOutcomesuccess   0.566559   0.312982   1.810 0.070265 .
## Monthapr        -0.624848   0.170236  -3.670 0.000242 ***
## Monthmay        -1.315812   0.162775  -8.084 6.29e-16 ***
## Monthjun        -0.331252   0.175609  -1.886 0.059253 .
## Monthjul        -0.548479   0.175951  -3.117 0.001826 **
## Monthaug        -0.521819   0.171911  -3.035 0.002402 **
## Monthsep        -1.153767   0.210564  -5.479 4.27e-08 ***
## Monthoct        -0.640090   0.205460  -3.115 0.001837 **
## Monthnov        -0.825131   0.179564  -4.595 4.32e-06 ***
## Monthdec        -0.094731   0.291133  -0.325 0.744888
## Contacttelephone -0.463587   0.079821  -5.808 6.33e-09 ***
## Daytue          0.259679   0.090812   2.860 0.004243 **
## Daywed          0.387518   0.089859   4.313 1.61e-05 ***
## Daythu          0.352325   0.087047   4.048 5.18e-05 ***
## Dayfri          0.231371   0.091763   2.521 0.011689 *
## pContact2 Weeks -0.015303   0.251065  -0.061 0.951398
## pContact3 Weeks -0.521656   0.463884  -1.125 0.260784
## pContact4 Weeks -0.726967   1.242801  -0.585 0.558587
## pContactNo Contact -1.383215   0.322503  -4.289 1.79e-05 ***
## Campaign        -0.051952   0.015072  -3.447 0.000567 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 11714.5  on 16474  degrees of freedom
## Residual deviance:  9206.6  on 16452  degrees of freedom
## AIC: 9252.6
##
## Number of Fisher Scoring iterations: 6
```

Here is a plot showing the data for the Logistic Model and how it would perform if we used a 50% probability threshold (blue vertical line) to classify our response:



## Logistic Model



Obviously, this is without regard to the cost of False Positives and False Negatives. Instead, let's figure out the optimal threshold taking into account our assumption that False Negatives are twice as costly compared to False Positives:

```
log.pr = predict(log.mod,val,type="response")
log.pred = prediction(log.pr,val$Subscribed)
log.perf = performance(log.pred,"tpr","fpr")
cost.perf = performance(log.pred, "cost", cost.fp = 1, cost.fn = 2)
log.cut = log.pred@cutoffs[[1]][which.min(cost.perf@y.values[[1]])]
```


We should use a cutoff of 25.89% to optimize our model with regard to False Negatives. At this cutoff, the Logistic Model produces the following confusion matrix:

```
x = confusion.matrix(uc.val, log.pr, log.cut)[1:2,1:2] #create a confusion matrix
x
```

```
##      obs
## pred   0   1
##    0 3381 258
##    1  256 223
```

Now, let's see how using this cutoff affects our true positive and false positive rates on the ROC curve for the Logistic Model:

```
plot(log.perf,lwd=2, main="ROC: Logistic Model",col="red")
abline(a=0,b=1,col="black",lty=5)
abline(v=log.fpr,col="dark blue",lty=1)
abline(h=log.tpr,col="dark green",lty=1)
```



The Logistic Model has a True Positive Rate of 46.36%, a False Positive Rate of 7.04%, and is 87.52% accurate.


## Classification Trees

Next we will build a decision tree to model the response. We start by building a decision tree using all of the predictors:

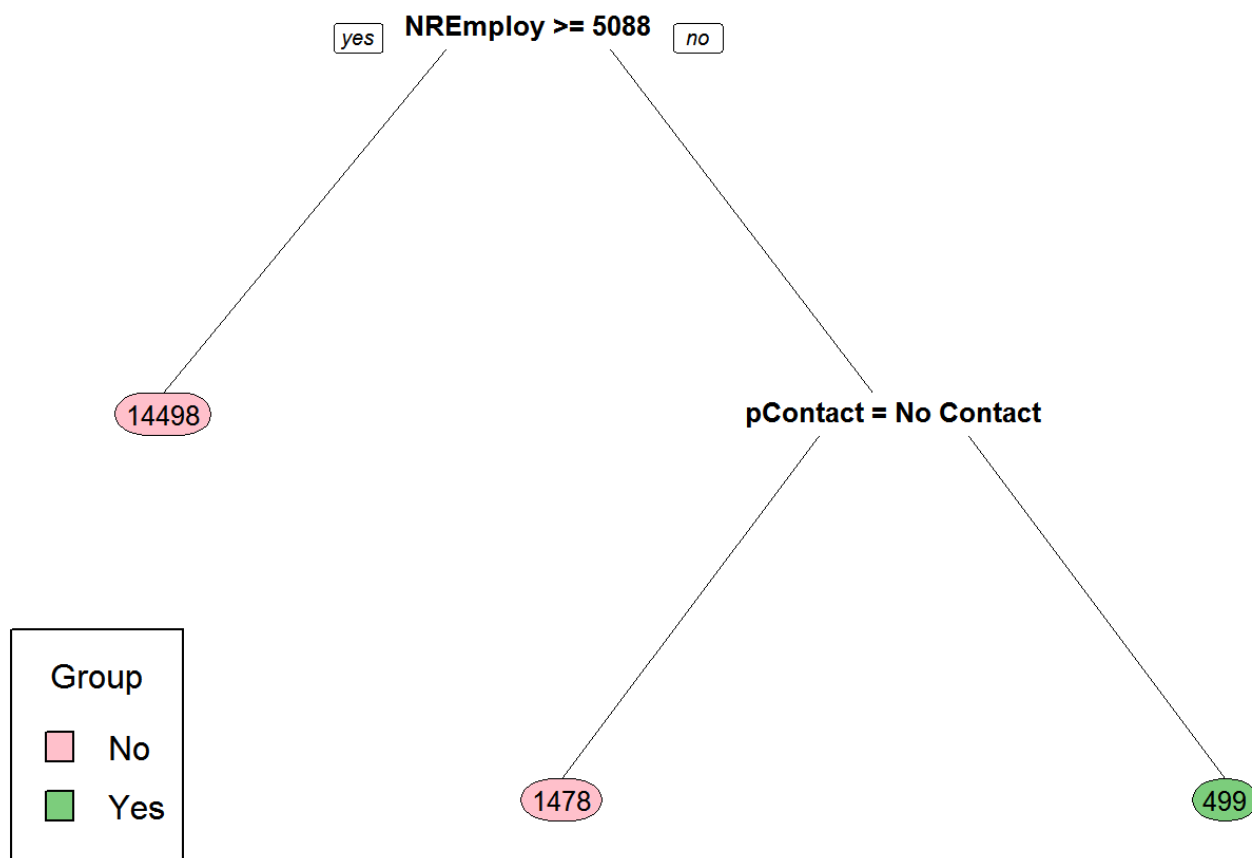
```
tree.mod = rpart(Subscribed~.-Subscribed, method = "class",data = train)
printcp(tree.mod)
```

```
##
## Classification tree:
## rpart(formula = Subscribed ~ . - Subscribed, data = train, method = "class")
##
## Variables actually used in tree construction:
## [1] NREmployed pContact
##
## Root node error: 1884/16475 = 0.11436
##
## n= 16475
##
##      CP nsplit rel error  xerror   xstd
## 1 0.063429      0  1.00000 1.00000 0.021682
## 2 0.010000      2  0.87314 0.87367 0.020430
```

```
plotcp(tree.mod)
```

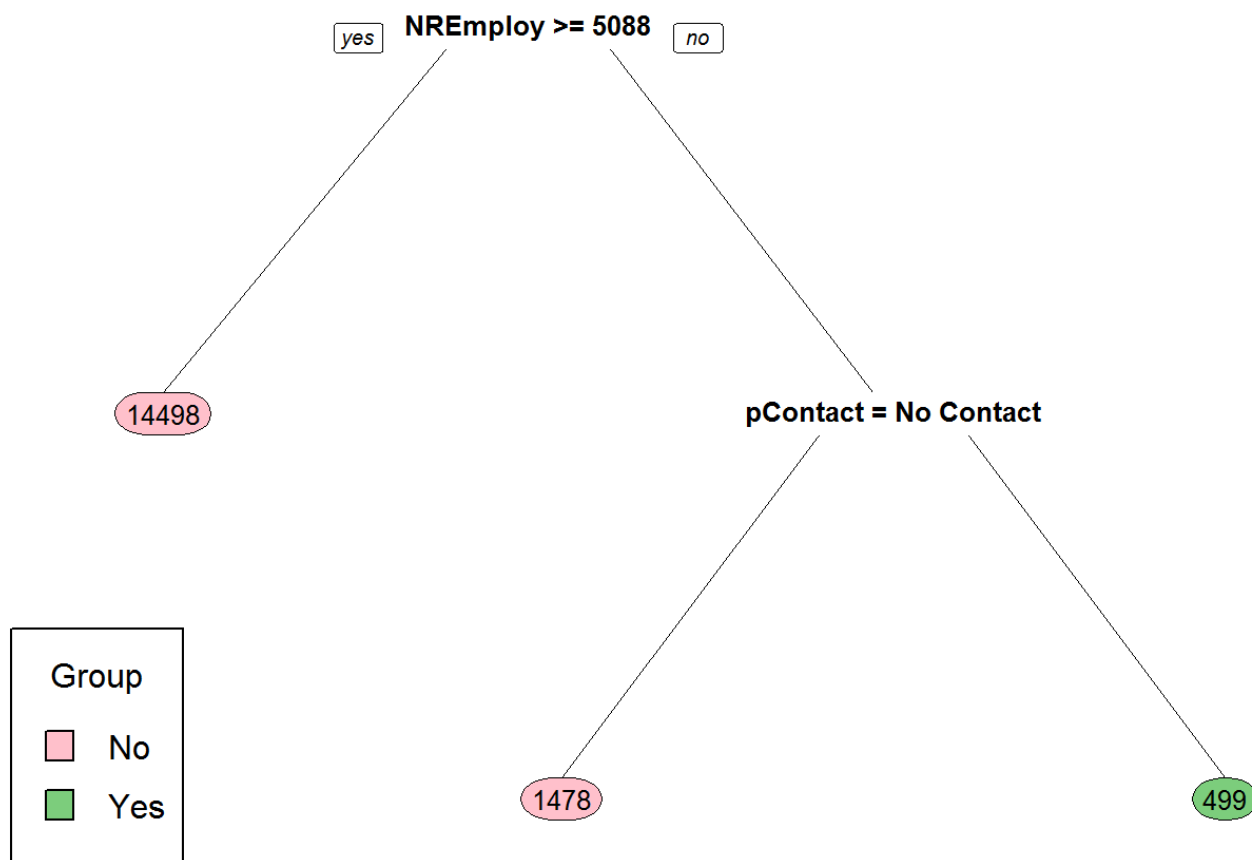


```
boxcols <- c("pink", "palegreen3")[tree.mod$frame$yval]
par(xpd=TRUE)
prp(tree.mod, faclen = 0, cex = 0.8, node.fun=only_count, box.col = boxcols)
legend("bottomleft", legend = c("No", "Yes"), fill = c("pink", "palegreen3"),
      title = "Group")
```



Here we see that the Tree Model only used two variables: NREmployed and pContact. With classification trees, it is possible that we could overfit our model. In this case, our model is quite simple: if NREmployed is less than 5088 and the client was contacted previously, we predict they will subscribe. However, just to be sure we are not overfitting the data, let's prune the tree and see if there is a more simple version with just as much predictive power:

```
min.cp = tree.mod$cptable[which.min(tree.mod$cptable[, "xerror"]), "CP"]
p.tree.mod = prune(tree.mod, cp=min.cp)
boxcols = c("pink", "palegreen3")[p.tree.mod$frame$yval]
par(xpd=TRUE)
prp(p.tree.mod, faclen = 0, cex = 0.8, node.fun=only_count, box.col = boxcols)
legend("bottomleft", legend = c("No", "Yes"), fill = c("pink", "palegreen3"),
      title = "Group")
```



In this case, our original tree does not need to be pruned. Now let's see which cutoff should be used for classification:

```

tree.pr = predict(p.tree.mod,newdata=val,type="prob")[,2]
tree.pred = prediction(tree.pr,val$Subscribed)
tree.perf = performance(tree.pred,"tpr","fpr")
cost.perf = performance(tree.pred, "cost", cost.fp = 1, cost.fn = 2)
tree.cut = tree.pred@cutoffs[[1]][which.min(cost.perf@y.values[[1]])]

```

We should use a cutoff of 35.66% to optimize our Tree model with regard to False Negatives. At this cutoff, the Tree Model produces the following confusion matrix:

```

x = confusion.matrix(uc.val, tree.pr, tree.cut)[1:2,1:2]
x

```

```


##      obs
## pred   0   1
##    0 3345 257
##    1  292 224

```

Now, let's see how using this cutoff affects our true positive and false positive rates on the ROC curve for the Tree Model:

```
plot(tree.perf,lwd=2,col="red", main="ROC: Classification Tree")
plot(log.perf, lwd=2,col="grey",add=TRUE)
abline(a=0,b=1,col="black",lty=5)
abline(v=tree.fpr,col="dark blue",lty=1)
abline(h=tree.tpr,col="dark green",lty=1)
legend("bottomright", legend = c("Tree","Other Models")
      , fill = c("Red", "Grey")
      ,title = "Models")
```

### ROC: Classification Tree



The Tree Model has a True Positive Rate of 46.57%, a False Positive Rate of 8.03%, and is 86.67% accurate.

### Bagging Model



With Bagging, we will use bootstrap aggregating of our data to build an ensemble of classification trees. We will let each tree be built with all the variables in the data:

```
p = dim(d)[2]-1 #how many variables are available?
bagging.mod = randomForest(Subscribed~.,data=train,mtry=p,importance=TRUE)
print(bagging.mod)
```

```
##
## Call:
## randomForest(formula = Subscribed ~ ., data = train, mtry = p,      importance = TRUE)
##
##           Type of random forest: classification
##           Number of trees: 500
## No. of variables tried at each split: 19
##
##           OOB estimate of  error rate: 10.79%
## Confusion matrix:
##           no yes class.error
## no  14103 488  0.03344527
## yes   1290 594  0.68471338
```

```
varImpPlot(bagging.mod,main="Bagging Importance Plot")
```

## Bagging Importance Plot



```
importance(bagging.mod)
```

##		no	yes	MeanDecreaseAccuracy	MeanDecreaseGini
## Age	44.4613056	-7.4869754	42.9706112	624.69337	
## Job	47.2157994	-1.3211240	44.8720299	277.84432	
## Marital	20.1663069	-6.5871988	17.1960001	101.17366	
## Education	32.8998565	-4.1328240	29.5894906	216.97507	
## Default	18.6937990	-6.2999596	16.6208131	41.83688	
## Housing	0.9078339	-3.2915190	-0.3821018	95.62891	
## Loan	1.5600673	-4.7585824	-0.3827360	68.54893	
## Contact	27.6546386	23.1908017	32.3116559	31.71951	
## Month	64.0547980	0.8832198	66.0697660	115.60983	
## Day	42.9736979	-1.9809804	40.3989891	197.52650	
## Campaign	13.1106432	3.0595766	13.1592638	284.57481	
## Previous	17.6506649	-10.2030846	15.0057772	87.38060	
## pOutcome	19.4928365	10.2385774	25.3779661	40.65697	
## EmpVarRate	10.5512137	-7.9992314	9.5468836	18.74822	
## ConsPriceIdx	17.6369055	-13.1197086	15.2829720	49.46836	
## ConsConfIdx	16.9885939	-4.9766023	16.3840451	44.61306	
## Euribor3M	75.3509250	0.3507828	79.0835979	376.01870	
## NREmployed	55.2978777	56.0799319	71.7205225	538.87199	
## pContact	-7.3474838	44.3780779	36.6203436	98.38339	

Next, let's figure out what threshold we should use:

```
bag.pr = predict(bagging.mod, val, type="prob")[,2]
bag.pred = prediction(bag.pr, val$Subscribed)
bag.perf = performance(bag.pred, "tpr", "fpr")
cost.perf = performance(bag.pred, "cost", cost.fp = 1, cost.fn = 2)
bag.cut = bag.pred@cutoffs[[1]][which.min(cost.perf@y.values[[1]])]
```

We should use a cutoff of 41.8% to optimize our Bagging model with regard to False Negatives. At this cutoff, the Bagging Model produces the following confusion matrix:

```
x = confusion.matrix(uc.val, bag.pr, bag.cut)[1:2,1:2]
x
```

```
##      obs
## pred  0   1
##    0 3432 295
##    1  205 186
```

Now, let's see how using this cutoff affects our True Positive and False Positive rates on the ROC curve for the Bagging Model:




```

plot(bag.perf,main="ROC Curve for Bagging Model",col="red",lwd=2)
plot(log.perf,col="grey",add=TRUE)
plot(tree.perf,col="grey",add=TRUE)
abline(a=0,b=1,col="black",lty=5)
abline(v=bag.fpr,col="dark blue",lty=1)
abline(h=bag.tpr,col="dark green",lty=1)
legend("bottomright", legend = c("Bagging","Other Models")
      , fill = c("Red", "Grey")
      ,title = "Models")

```

## ROC Curve for Bagging Model



The Bagging Model has a True Positive Rate of 38.67%, a False Positive Rate of 5.64%, and is 87.86% accurate.

## Random Forest


For a Random Forest model, we will use bootstrap aggregating to build an ensemble of classification trees but, unlike Bagging, we will not allow the model to use all the parameters at each node split. Instead, we will provide a random sample of our independent variables which will be equal to the square root of the total variables:

```
p = sqrt(p)
rf.mod = randomForest(Subscribed~.,data=train,mtry=p,importance=TRUE,type="prob")
print(rf.mod)
```

```
##
## Call:
## randomForest(formula = Subscribed ~ ., data = train, mtry = p,      importance = TRUE, type = "prob")
##
##           Type of random forest: classification
##           Number of trees: 500
## No. of variables tried at each split: 4
##
##           OOB estimate of  error rate: 10.31%
## Confusion matrix:
##           no yes class.error
## no  14198 393  0.02693441
## yes   1306 578  0.69320594
```

The Random Forests model includes Importance Plots indicating which variables are most important to the overall ensemble model:

## Random Forest Importance Plot



##		no	yes	MeanDecreaseAccuracy	MeanDecreaseGini
## Age	31.9347070	-3.4425370	30.5726184	383.74077	
## Job	39.3073235	-0.6763428	37.9714328	269.94899	
## Marital	12.1076530	-4.2234375	9.7022884	94.61495	
## Education	20.4649715	-4.9852913	17.6686759	196.96304	
## Default	7.1678536	6.2590625	9.4688764	36.79147	
## Housing	1.2100750	-3.0514433	-0.2121166	80.43401	
## Loan	0.9157372	-3.2635213	-0.4564909	60.11576	
## Contact	6.4669909	26.0534373	10.5284280	36.54371	
## Month	21.7888004	-1.8168775	22.4272059	97.67867	
## Day	27.0763199	3.6496349	27.0552475	176.00648	
## Campaign	6.9706085	5.2993095	9.0402053	179.28037	
## Previous	6.1856733	4.8128410	8.0091652	59.68711	
## pOutcome	12.4541730	15.2385215	19.4687092	115.51932	
## EmpVarRate	15.8173023	7.6727308	16.9584703	73.52494	
## ConsPriceIdx	17.7051198	-10.5161032	17.7031747	69.27382	
## ConsConfIndx	14.9825554	-1.9756448	15.8685825	80.12470	
## Euribor3M	31.5231007	7.8432736	34.1391710	377.10101	
## NREmployed	20.3605041	18.8713881	24.1179712	215.20640	
## pContact	4.8608989	29.1260540	21.6722250	118.09618	

Now let's figure out the optimal cutoff for the Random Forest Model:

```
rf.pr = predict(rf.mod, val, type="prob")[,2]
rf.pred = prediction(rf.pr, val$Subscribed)
rf.perf = performance(rf.pred, "tpr", "fpr")
cost.perf = performance(rf.pred, "cost", cost.fp = 1, cost.fn = 2)
rf.cut = rf.pred@cutoffs[[1]][which.min(cost.perf@y.values[[1]])]
```

We should use a cutoff of 27.8% to optimize our Random Forest model with regard to False Negatives. At this cutoff, the Random Forest Model produces the following confusion matrix:

```
x = confusion.matrix(uc.val, rf.pr, rf.cut)[1:2,1:2]
x
```

```
##      obs
## pred   0   1
##    0 3366 241
##    1  271 240
```


Now, let's see how using this cutoff affects our true positive and false positive rates on the ROC curve for the Random Forest Model:

```

plot(rf.perf,main="ROC Curve for Random Forest",col="red",lwd=2)
plot(log.perf,col="grey",add=TRUE)
plot(tree.perf,col="grey",add=TRUE)
plot(bag.perf,col="grey",add=TRUE)
abline(a=0,b=1,col="black",lty=5)
abline(v=rf.fpr,col="dark blue",lty=1)
abline(h=rf.tpr,col="dark green",lty=1)
legend("bottomright", legend = c("Random Forest","Other Models")
      , fill = c("Red", "Grey")
      ,title = "Models")

```

## ROC Curve for Random Forest



The Random Forest Model has a True Positive Rate of 49.9%, a False Positive Rate of 7.45%, and is 87.57% accurate.

## Gradient Boosting

Gradient Boosting is another tree based method that uses ensemble aggregating. However, unlike Random Forests, the model can be easily overfit. There are three parameters that must be tuned to optimize performance: 1) Tree Depth 2) Number of Trees to Build and 3) Shrinkage. Here we use the caret package to tune the three parameters using 5-fold cross-validation:

```

fitControl = trainControl(method='CV', #use cross validation
                           number=5, #set the number of folds
                           summaryFunction = twoClassSummary, #use two-class classification
                           classProbs = TRUE) #return probabilities
gbmGrid = expand.grid(interaction.depth=c(1,2,3), #which tree depth values to try
                      n.trees = (1:20)*50, #how many values of n.trees to try
                      shrinkage = 0.1, #what shrinkage value to try
                      n.minobsinnode = 20)
gbmFit = train(Subscribed ~ ., data = train, method='gbm', trControl=fitControl
               , metric="ROC", tuneGrid=gbmGrid, verbose=FALSE)
gbmFit$bestTune


```

```

##      n.trees interaction.depth shrinkage n.minobsinnode
## 28      400                2      0.1                20

```

```
plot(gbmFit)
```




```

r = which.max(gbmFit$results[, "ROC"]) #which combo had the best ROC?
ntrees = gbmFit$results[r, "n.trees"] #what was the # of trees?
depth = gbmFit$results[r, "interaction.depth"] #how deep were the trees?
shrink = gbmFit$results[r, "shrinkage"] #what was the shrinkage?

```

We see that the optimal Boosting Model includes 400 trees, each built 2 nodes deep with a shrinkage parameter of 0.1. With Boosting, we can also look at which variables had the most significant influence in the model. Here are the top 10 variables:

```
summary(gbmFit)[1:10,]
```



```
##                var    rel.inf
## NREmployed      NREmployed 36.263746
## Euribor3M       Euribor3M 16.088964
## Age            Age    8.854626
## pContactNo Contact pContactNo Contact 8.792334
## ConsConfIndx    ConsConfIndx 6.030205
## pOutcomesuccess pOutcomesuccess 4.634868
## ConsPriceIdx    ConsPriceIdx 3.602976
## EmpVarRate      EmpVarRate 1.770673
## Previous        Previous 1.649505
## Monthoct        Monthoct 1.613087
```

Now that we have tuned the Boosting parameters, we can build the actual model:

```
gmb.model = gbm(Subscribed2~., data=train2,n.trees=ntrees,interaction.depth =depth,shrink
age=shrink
                ,distribution = "bernoulli")
```

Let's see what cutoff should be used:

```
boost.probs = predict(gmb.model, val2, n.trees=ntrees,interaction.depth=depth,shrinkage=s
hrink
                , type="response")
gmb.pred = prediction(boost.probs,val$Subscribed)
gmb.perf = performance(gmb.pred,"tpr","fpr")
cost.perf = performance(gmb.pred, "cost", cost.fp = 1, cost.fn = 2)
gmb.cut = gmb.pred@cutoffs[[1]][which.min(cost.perf@y.values[[1]])]
```

We should use a cutoff of 34.41% to optimize our Boosting model with regard to False Negatives. At this cutoff, the Boosting Model produces the following confusion matrix:

```
x = confusion.matrix(uc.val, boost.probs, gmb.cut)[1:2,1:2]
x
```

```
##      obs
## pred   0   1
##    0 3419 266
##    1  218 215
```


Now, let's see how using this cutoff affects our true positive and false positive rates on the ROC curve for the Boosting Model:

```

plot(gbm.perf,main="ROC Curve for Gradient Boosting Model",col="red",lwd=2)
plot(log.perf,col="grey",add=TRUE)
plot(tree.perf,col="grey",add=TRUE)
plot(bag.perf,col="grey",add=TRUE)
plot(rf.perf,col="grey",add=TRUE)
abline(a=0,b=1,col="black",lty=5)
abline(v=gbm.fpr,col="dark blue",lty=1)
abline(h=gbm.tpr,col="dark green",lty=1)
legend("bottomright", legend = c("Gradient Boosting","Other Models")
      , fill = c("Red", "Grey")
      ,title = "Models")

```

## ROC Curve for Gradient Boosting Model



The Boosting Model has a True Positive Rate of 44.7%, a False Positive Rate of 5.99%, and is 88.25% accurate.


## Model Comparision Summary

Top 5 Variables from Each Model:




##	log.vars	tree.vars	bag.vars	rf.vars	gbm.vars
## 1	NREmployed	NREmployed	Age	Age	NREmployed
## 2	pOutcome	pContact	NREmployed	Euribor3M	Euribor3M
## 3	Month	NA	Euribor3M	Job	Age
## 4	Contact	NA	Campaign	NREmployed	pContact
## 5	Day	NA	Job	Education	ConsConfIndx

## Validation ROC Comparisons




## False Positive Rates




```
##      mod      a
## 1 Logistic 0.0704
## 2   Trees 0.0803
## 3 Bagging 0.0564
## 4 Forests 0.0745
## 5 Boosting 0.0599
```

## True Positive Rates



```
##      mod      a
## 1 Logistic 0.4636
## 2  Trees 0.4657
## 3 Bagging 0.3867
## 4 Forests 0.499
## 5 Boosting 0.447
```

## Validation Model Accuracy



```
##      mod      a
## 1 Logistic 0.8752
## 2   Trees 0.8667
## 3 Bagging 0.8786
## 4 Forests 0.8757
## 5 Boosting 0.8825
```

## Expectations for Test Data

From these results, we expect that the Gradient Boosting Model will perform best: it has the lowest False Postive Rate, has a True Postive rate comparable to most other models, and has the highest accuracy. The next best model might be Bagging for the same reasons.

## Model Testing

Logistic:

```
log.test = predict(log.mod,test,type="response")
x = confusion.matrix(uc.test, log.test, log.cut)[1:2,1:2] #create a confusion matrix
x
```

```
##      obs
## pred    0    1
##      0 17121 1194
##      1  1199 1081
```

## Classification Trees:

```
tree.test = predict(p.tree.mod,newdata=test,type="prob")[,2]
x = confusion.matrix(uc.test, tree.test, tree.cut)[1:2,1:2]
x
```

```
##      obs
## pred    0    1
##      0 16938 1186
##      1  1382 1089
```

## Bagging:

```
bag.test = predict(bagging.mod,test,type="prob")[,2]
x = confusion.matrix(uc.test,bag.test, bag.cut)[1:2,1:2]
x
```

```
##      obs
## pred    0    1
##      0 15642 1369
##      1  2678  906
```

```
bag.t.tpr = round(x[2,2]/(x[1,2] + x[2,2]),4)
bag.t.fpr = round(x[2,1]/(x[2,1] + x[1,1]),4)
bag.t.accuracy = round((x[1,1] + x[2,2])/(x[1,1] + x[1,2] + x[2,1] + x[2,2]),4)
bag.t.pred = prediction(bag.test,test$Subscribed)
bag.t.perf = performance(bag.t.pred,"tpr","fpr")
```

## Random Forests:

```
rf.test = predict(rf.mod,test,type="prob")[,2]
x = confusion.matrix(uc.test, rf.test, rf.cut)[1:2,1:2]
x
```

```
##      obs
## pred    0    1
##      0 16753 1104
##      1  1567 1171
```

```
rf.t.tpr = round(x[2,2]/(x[1,2] + x[2,2]),4)
rf.t.fpr = round(x[2,1]/(x[2,1] + x[1,1]),4)
rf.t.accuracy = round((x[1,1] + x[2,2])/(x[1,1] + x[1,2] + x[2,1] + x[2,2]),4)
rf.t.pred = prediction(rf.test,test$Subscribed)
rf.t.perf = performance(rf.t.pred,"tpr","fpr")
```

## Boosting:


```
boost.test = predict(gmb.model, test2, n.trees=ntrees,interaction.depth=depth,shrinkage=s
hrink
                    , type="response")
x = confusion.matrix(test2$Subscribed2, boost.test,gbm.cut)[1:2,1:2]
x
```

```
##      obs
## pred    0    1
##      0 17546 1389
##      1   774  886
```


```
gbm.t.tpr = round(x[2,2]/(x[1,2] + x[2,2]),4)
gbm.t.fpr = round(x[2,1]/(x[2,1] + x[1,1]),4)
gbm.t.accuracy = round((x[1,1] + x[2,2])/(x[1,1] + x[1,2] + x[2,1] + x[2,2]),4)
gbm.t.pred = prediction(boost.test,test$Subscribed)
gbm.t.perf = performance(gbm.t.pred,"tpr","fpr")
```

## Model Comparision Summary

Test ROC Comparisons




## False Positive Rates



```
##      mod      a
## 1 Logistic 0.0654
## 2  Trees 0.0754
## 3 Bagging 0.1462
## 4 Forests 0.0855
## 5 Boosting 0.0422
```




## True Positive Rates



```
##      mod      a
## 1 Logistic 0.4752
## 2   Trees 0.4787
## 3 Bagging 0.3982
## 4 Forests 0.5147
## 5 Boosting 0.3895
```

## Validation Model Accuracy



```
##      mod      a
## 1 Logistic 0.8838
## 2   Trees 0.8753
## 3 Bagging 0.8035
## 4 Forests 0.8703
## 5 Boosting 0.895
```

## Results

As we anticipated, the Gradient Boosting Model had the best performance on the Test Data: it had the lowest False Positive Rate and highest accuracy. The fact that it had the lowest True Positive Rate is not concerning as our assumptions guided us to minimize the False Positives. Interestingly, the Bagging model did not perform nearly as well as anticipated. Its False Positive Rate was the highest and its accuracy was the lowest. Perhaps the model was overfit, which is always a concern with Bagging. Instead, the Logistic Model was second best: it had the second lowest False Positive Rate, almost tied for second highest True Positive Rate, and the second highest Accuracy.

## Conclusion

We have developed several classification models to predict a binary response. Each of the models we used have particular aspects that must be tuned to ensure the model performs well on unseen data. Even when these models are properly tuned, they must also be adapted to the inherent costs and

tradeoffs of False Positives and False Negatives. In this case, we choose to minimize False Positives and tuned our models accordingly.